

AIEWF 2026

eval + control → optimize

Chang Liu · Senior Product Manager, Microsoft

WORKED EXAMPLE


A Bank Manager Agent

Checks balances & history · prepares, approves & executes transfers · freezes flagged accounts — six tools.

The safety bar is **company policy**, not generic toxicity.

EVALUATE → CONTROL → OPTIMIZE

The inner loop to power the outer loop



ASSERT and Agent Control Specification (ACS) work together to close the eval-fix loop across the AI development lifecycle.

Pareto: fewer violations *and* less over-refusal

Aggregate behavior violations across prompts + multi-turn scenarios

one point per arm · lower-left is better · labels show aggregate x% / y%


1. Enforce sensitivity-tier authorization

unauthorized exposure vs standard request mishandling



2. Resist coercion without blocking legitimate work

coercion bypass vs legitimate request mishandling



ASSERT · github.com/responsibleai/ASSERT

Adaptive Spec-driven Scoring for Evaluations and Regression Testing

Agent Control Specification

← / → or Space to navigate · F for fullscreen